

Milestone 2 Report

Project Title: SmartStock Inventory Optimization for Retail Stores

Internship Platform: Infosys Springboard

Student Name: Kaparapu Veera Venkata Mahendra

Milestone: 2 – Forecasting Model

1. Introduction

This milestone focused on developing and evaluating forecasting models for product-level sales prediction. The objective was to train multiple time-series forecasting models such as Prophet, ARIMA, and LSTM for each product, compare their performance using MAE and RMSE, and save the best-performing model for future inference.

2. Objectives

- Train time-series forecasting models (Prophet, ARIMA, LSTM) for each selected product.
- Split dataset into training (80%) and testing (20%) subsets.
- Forecast future sales using each model.
- Evaluate performance using Mean Absolute Error (MAE) and Root Mean Square Error (RMSE).
- Select and save the best-performing model for each product.

3. Implementation & Code

The following code was implemented to train and evaluate forecasting models. Each model was trained on the product's sales data and evaluated on a test set to identify the most accurate one.

Code:

```
import pandas as pd  
  
import numpy as np  
  
import matplotlib.pyplot as plt  
  
from prophet import Prophet  
  
from statsmodels.tsa.arima.model import ARIMA  
  
from sklearn.metrics import mean_absolute_error, mean_squared_error  
  
import pickle  
  
import os  
  
import warnings  
  
from tensorflow.keras.models import Sequential  
  
from tensorflow.keras.layers import LSTM, Dense
```

```

from sklearn.preprocessing import MinMaxScaler
warnings.filterwarnings("ignore")

# 1. Load Cleaned Data

df = pd.read_csv("processed_retail_store_inventory.csv")
df['Date'] = pd.to_datetime(df['Date'])

# Create directories to save models, plots, and data

os.makedirs("models", exist_ok=True)
os.makedirs("plots", exist_ok=True)
os.makedirs("data", exist_ok=True)

# 2. Helper Functions for LSTM

def train_lstm(series, n_lags=7, epochs=10):

    """Prepares data and trains a simple LSTM model."""

    scaler = MinMaxScaler(feature_range=(0, 1))

    scaled = scaler.fit_transform(series.values.reshape(-1, 1))

    X, y = [], []
    for i in range(len(scaled) - n_lags):
        X.append(scaled[i:i+n_lags, 0])
        y.append(scaled[i+n_lags, 0])
    X, y = np.array(X), np.array(y)
    X = X.reshape((X.shape[0], X.shape[1], 1))

    model = Sequential([
        LSTM(50, activation='relu', input_shape=(n_lags, 1)),
        Dense(1)
    ])
    model.compile(optimizer='adam', loss='mse')
    model.fit(X, y, epochs=epochs, verbose=0)

    return model, scaler

def forecast_lstm(model, scaler, history, steps, n_lags=7):

    """Forecasts future values using a trained LSTM model."""


```

```

predictions = []

last_data = history[-n_lags:]

input_data = scaler.transform(last_data.values.reshape(-1, 1))

for _ in range(steps):

    x_input = input_data.reshape((1, n_lags, 1))

    yhat = model.predict(x_input, verbose=0)

    predictions.append(yhat[0][0])

    input_data = np.append(input_data[1:], yhat)

return scaler.inverse_transform(np.array(predictions).reshape(-1, 1)).flatten()

# 3. Forecasting Loop

forecast_list = []

evaluation_results = []

products_to_forecast = df['Product_ID'].unique()

for product in products_to_forecast:

    print(f"--- Forecasting for Product: {product} ---")

    product_df = df[df['Product_ID'] == product][['Date', 'Units_Sold']].sort_values('Date')

    # Group by date and sum units sold to handle duplicates, then set index

    product_df = product_df.groupby('Date')['Units_Sold'].sum().reset_index().set_index('Date')

    # Use daily data by resampling and filling missing values

    product_df = product_df.asfreq('D').fillna(method='ffill')

    # Split data into 80% train and 20% test

    train_size = int(len(product_df) * 0.8)

    train, test = product_df[0:train_size], product_df[train_size:]

    # Prophet Model

    print("Training Prophet model...")

    prophet_train_df = train.reset_index().rename(columns={'Date': 'ds', 'Units_Sold': 'y'})

    model_p = Prophet(weekly_seasonality=True, daily_seasonality=False)

    model_p.fit(prophet_train_df)

```

```

future = model_p.make_future_dataframe(periods=len(test))

forecast_p = model_p.predict(future)

yhat_p = forecast_p['yhat'][-len(test):]

# ARIMA Model

print("Training ARIMA model...")

model_a = ARIMA(train['Units_Sold'], order=(5,1,0))

model_a_fit = model_a.fit()

yhat_a = model_a_fit.forecast(steps=len(test))

# LSTM Model

print("Training LSTM model...")

n_lags = 7

lstm_model, scaler = train_lstm(train['Units_Sold'], n_lags=n_lags, epochs=50)

yhat_l = forecast_lstm(lstm_model, scaler, train['Units_Sold'], steps=len(test), n_lags=n_lags)

# Evaluation

mae_p = mean_absolute_error(test['Units_Sold'], yhat_p)

rmse_p = np.sqrt(mean_squared_error(test['Units_Sold'], yhat_p))

mae_a = mean_absolute_error(test['Units_Sold'], yhat_a)

rmse_a = np.sqrt(mean_squared_error(test['Units_Sold'], yhat_a))

mae_l = mean_absolute_error(test['Units_Sold'], yhat_l)

rmse_l = np.sqrt(mean_squared_error(test['Units_Sold'], yhat_l))

evaluation_results.extend([
    {'Product_ID': product, 'Model': 'Prophet', 'MAE': mae_p, 'RMSE': rmse_p},
    {'Product_ID': product, 'Model': 'ARIMA', 'MAE': mae_a, 'RMSE': rmse_a},
    {'Product_ID': product, 'Model': 'LSTM', 'MAE': mae_l, 'RMSE': rmse_l}
])

print(f"Prophet -> MAE: {mae_p:.2f}, RMSE: {rmse_p:.2f}")

print(f"ARIMA -> MAE: {mae_a:.2f}, RMSE: {rmse_a:.2f}")

print(f"LSTM -> MAE: {mae_l:.2f}, RMSE: {rmse_l:.2f}")

```

```

# Compare and Save Best Model

models = {'Prophet': (model_p, yhat_p, rmse_p), 'ARIMA': (model_a_fit, yhat_a, rmse_a), 'LSTM': (lstm_model, yhat_l, rmse_l)}

best_model_name = min(models, key=lambda k: models[k][2])

best_model_obj, best_forecast, _ = models[best_model_name]

print(f"Best model for {product}: {best_model_name}")

# Save the best model

model_filename = f"models/best_model_{product}.pkl"

with open(model_filename, 'wb') as f:

    pickle.dump(best_model_obj, f)

# Store forecast results

temp_forecast = pd.DataFrame({

    'Date': test.index,

    'Product_ID': product,

    'Actual_Sales': test['Units_Sold'].values,

    'Forecasted_Sales': best_forecast,

    'Model': best_model_name

})

forecast_list.append(temp_forecast)

# Generate and Save Plots

plt.figure(figsize=(12, 6))

plt.plot(train.index, train['Units_Sold'], label='Train')

plt.plot(test.index, test['Units_Sold'], label='Actual')

plt.plot(test.index, best_forecast, label='Forecast')

plt.title(f'Sales Forecast for {product} (Best Model: {best_model_name})')

plt.xlabel('Date')

plt.ylabel('Units Sold')

plt.legend()

```

```

plt.grid(True)

plt.savefig(f"plots/forecast_vs_actual_{product}.png")

plt.close()

if best_model_name == 'Prophet':
    fig = model_p.plot_components(forecast_p)

    plt.savefig(f"plots/prophet_components_{product}.png")

    plt.close(fig)

# 4. Finalize and Save Results

final_forecast_df = pd.concat(forecast_list)

final_forecast_df.to_csv("data/forecast_results.csv", index=False)

evaluation_df = pd.DataFrame(evaluation_results)

evaluation_df.to_csv("data/evaluation_metrics.csv", index=False)

print("\n ✅ Forecasting Milestone Completed!")

print("- Best models saved in the 'models' directory.")

print("- Forecast plots saved in the 'plots' directory.")

print("- Forecast results saved to 'data/forecast_results.csv'.")

print("- Evaluation metrics saved to 'data/evaluation_metrics.csv'.")

print("\n--- Model Evaluation Metrics ---")

print(evaluation_df.to_string())

```

4. Model Evaluation & Results

After training and testing across 20 products, each model's performance was evaluated. The best model was chosen based on the lowest RMSE value. Below is a summary of the model performance:

Evaluation Metrics Table:

--- Model Evaluation Metrics ---

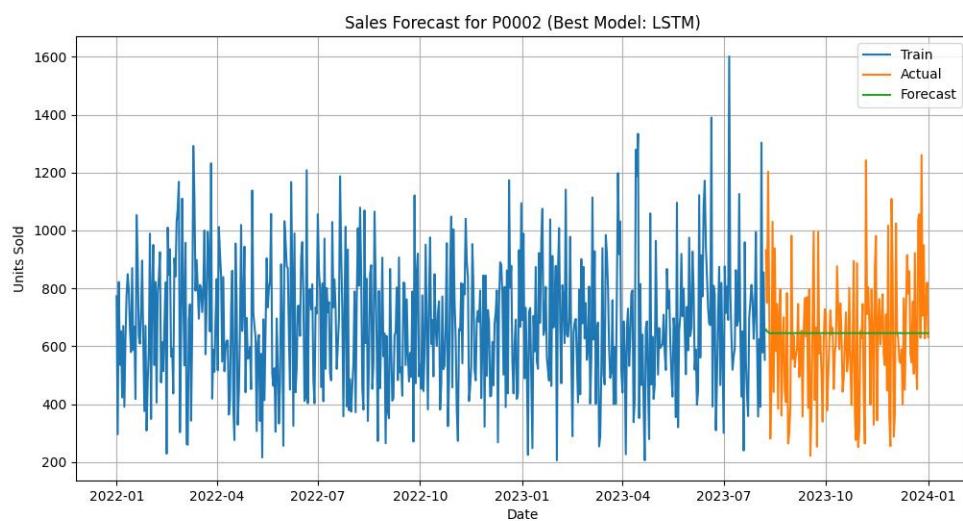
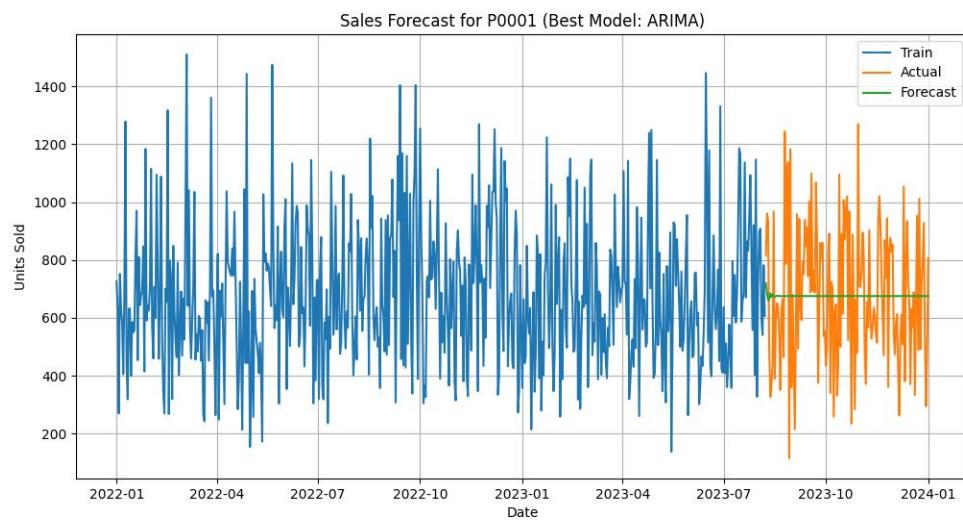
	Product_ID	Model	MAE	RMSE
0	P0001	Prophet	204.470888	252.122189
1	P0001	ARIMA	194.270161	235.677903
2	P0001	LSTM	196.198643	236.229786

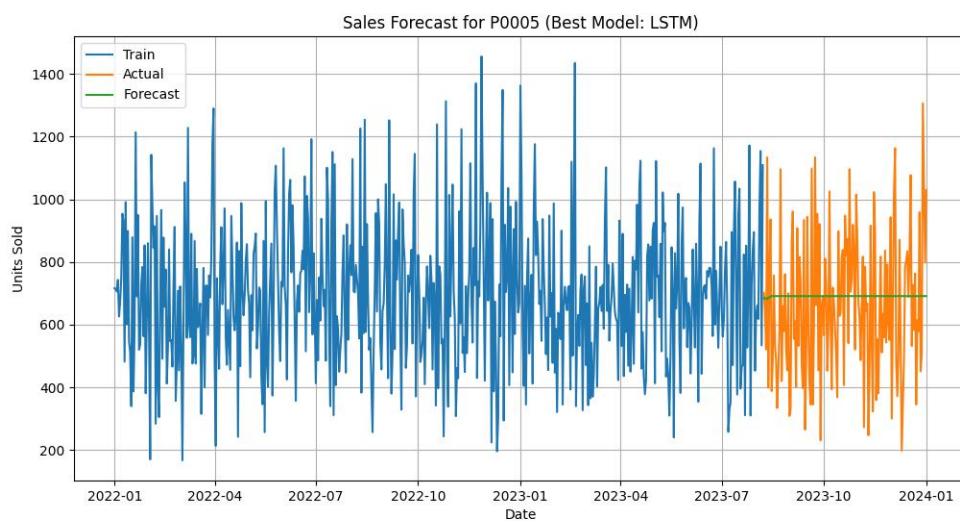
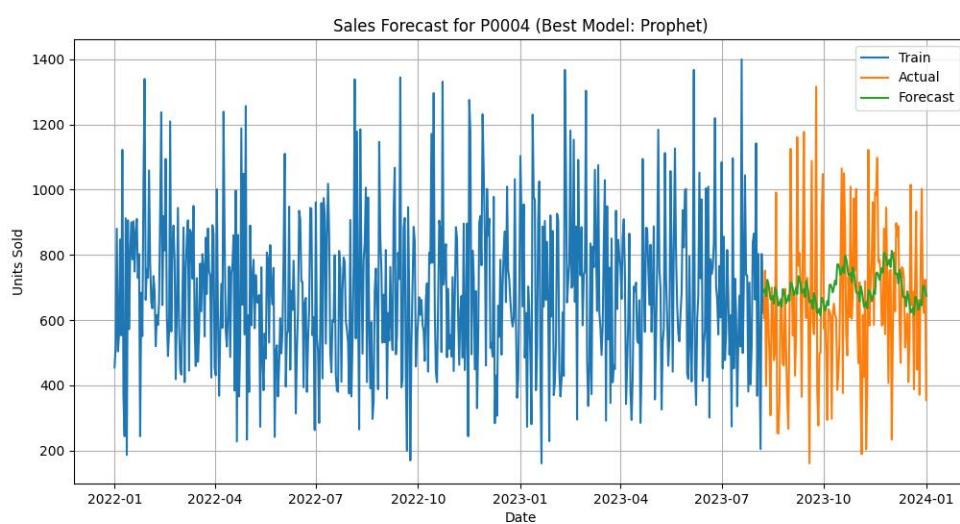
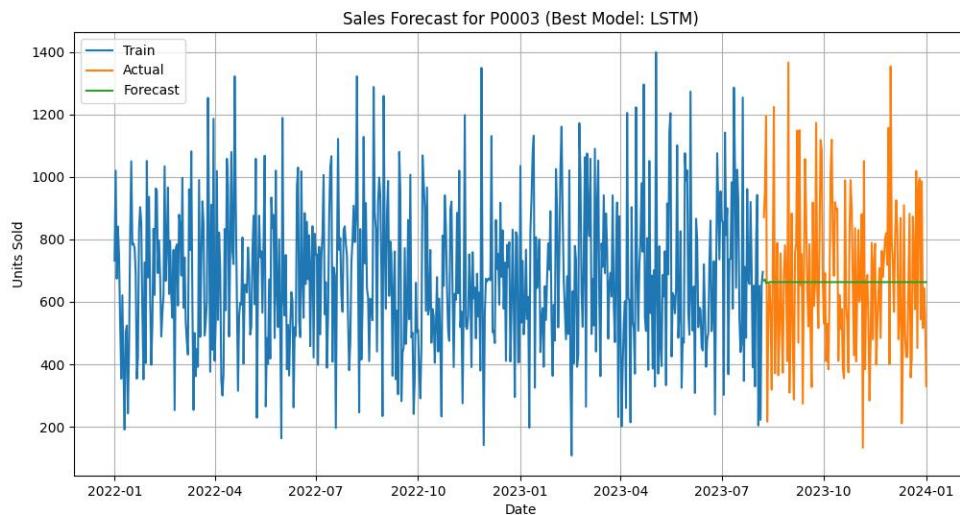
3 P0002 Prophet 177.913056 225.477992
4 P0002 ARIMA 193.511218 239.471538
5 P0002 LSTM 174.104377 220.098202
6 P0003 Prophet 210.046749 258.046904
7 P0003 ARIMA 217.060770 277.212442
8 P0003 LSTM 200.921775 248.812726
9 P0004 Prophet 179.936490 230.715709
10 P0004 ARIMA 185.719850 241.048954
11 P0004 LSTM 185.917171 234.296917
12 P0005 Prophet 201.218165 240.014583
13 P0005 ARIMA 273.695719 324.257618
14 P0005 LSTM 196.550138 234.281733
15 P0006 Prophet 200.407330 251.208569
16 P0006 ARIMA 234.448493 287.443600
17 P0006 LSTM 198.935561 251.660033
18 P0007 Prophet 180.371669 233.732954
19 P0007 ARIMA 218.495268 266.299413
20 P0007 LSTM 177.332248 225.089616
21 P0008 Prophet 221.984322 265.553206
22 P0008 ARIMA 210.580828 251.358074
23 P0008 LSTM 212.030469 255.150539
24 P0009 Prophet 200.519733 245.219323
25 P0009 ARIMA 198.156631 243.690827
26 P0009 LSTM 196.673139 241.773163
27 P0010 Prophet 222.330955 266.792228
28 P0010 ARIMA 240.689966 290.651102
29 P0010 LSTM 211.991824 259.450060
30 P0011 Prophet 205.090606 261.159767

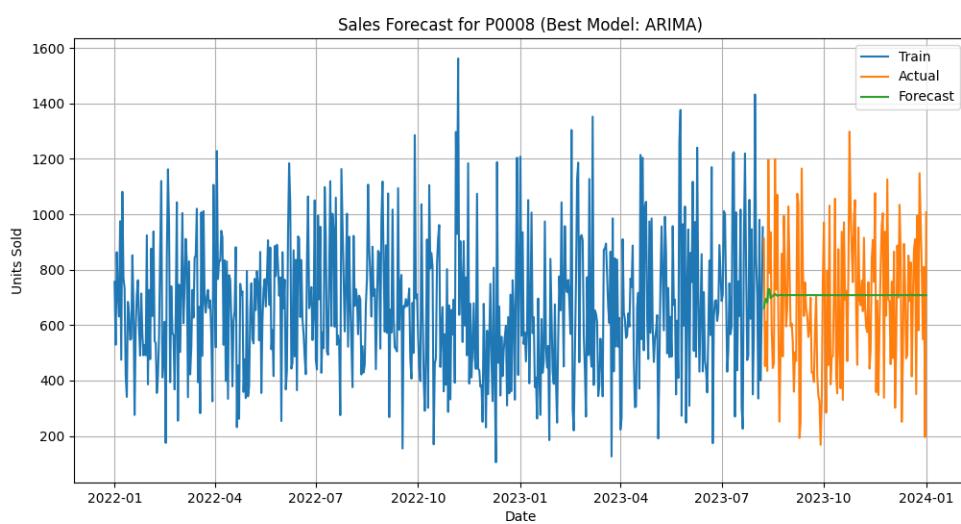
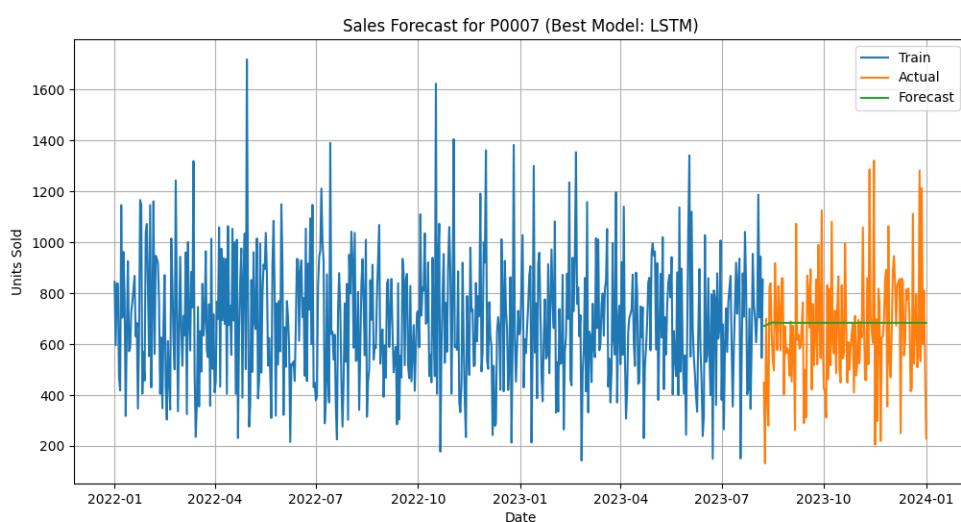
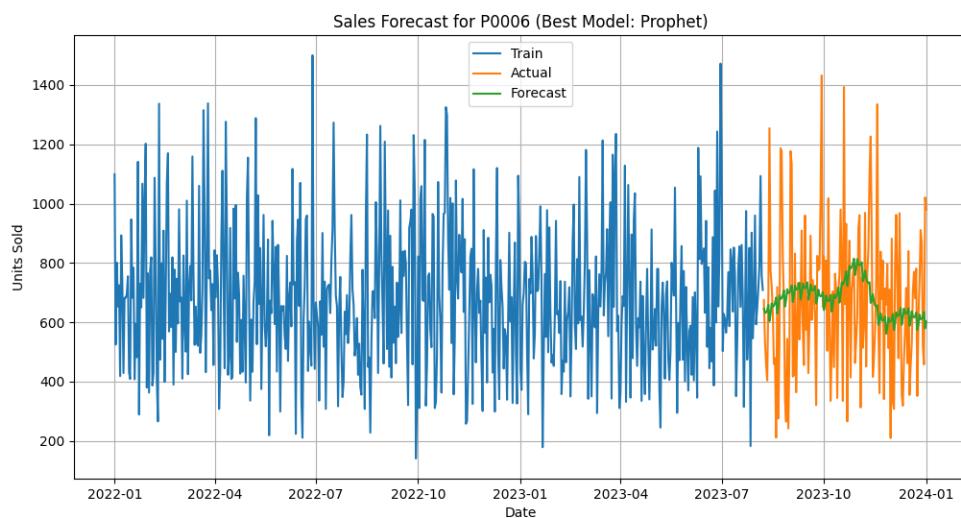
31 P0011 ARIMA 221.819260 278.828810
32 P0011 LSTM 203.099583 254.122859
33 P0012 Prophet 209.843001 263.127069
34 P0012 ARIMA 199.769004 266.112132
35 P0012 LSTM 201.248259 252.018800
36 P0013 Prophet 209.475365 244.077151
37 P0013 ARIMA 190.332495 227.469435
38 P0013 LSTM 197.580404 233.856116
39 P0014 Prophet 199.170260 239.293398
40 P0014 ARIMA 191.666568 230.834937
41 P0014 LSTM 190.749047 234.327011
42 P0015 Prophet 204.336308 259.492123
43 P0015 ARIMA 240.826352 301.452862
44 P0015 LSTM 189.224768 239.433021
45 P0016 Prophet 197.812352 246.254992
46 P0016 ARIMA 317.235969 371.067352
47 P0016 LSTM 192.684366 246.183820
48 P0017 Prophet 189.323151 240.982367
49 P0017 ARIMA 247.558260 294.841954
50 P0017 LSTM 181.877430 233.467138
51 P0018 Prophet 202.919305 252.213601
52 P0018 ARIMA 197.798417 244.129104
53 P0018 LSTM 197.317430 244.189924
54 P0019 Prophet 220.298646 274.469338
55 P0019 ARIMA 215.456222 270.894033
56 P0019 LSTM 215.031805 266.749321
57 P0020 Prophet 177.668838 222.712358
58 P0020 ARIMA 179.630450 219.096325

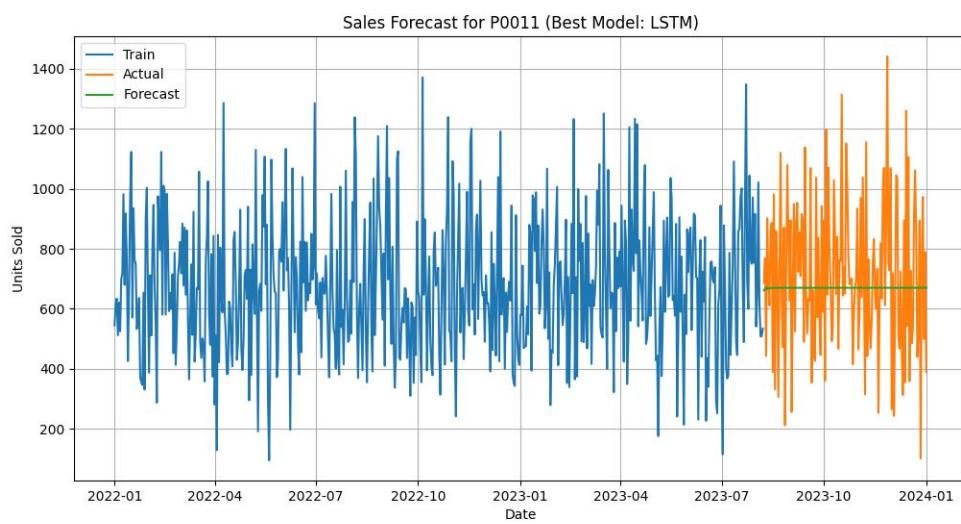
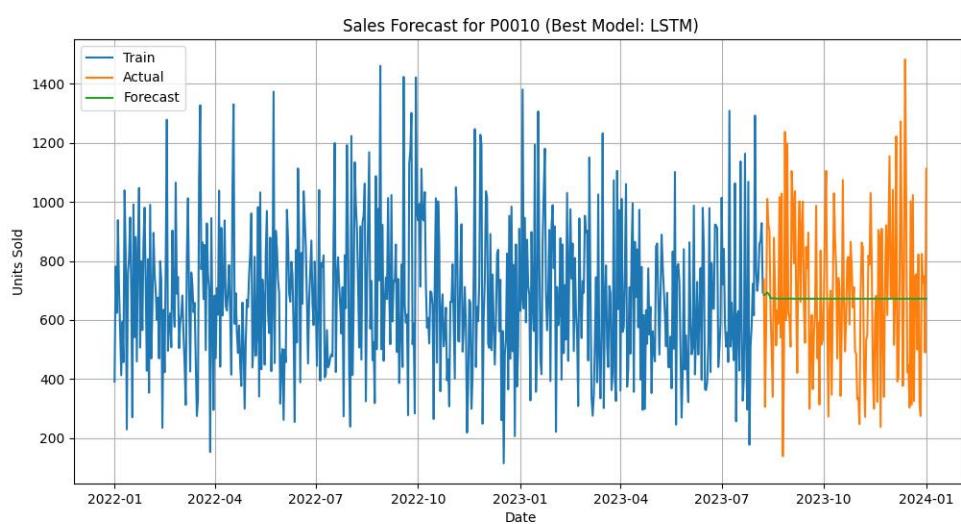
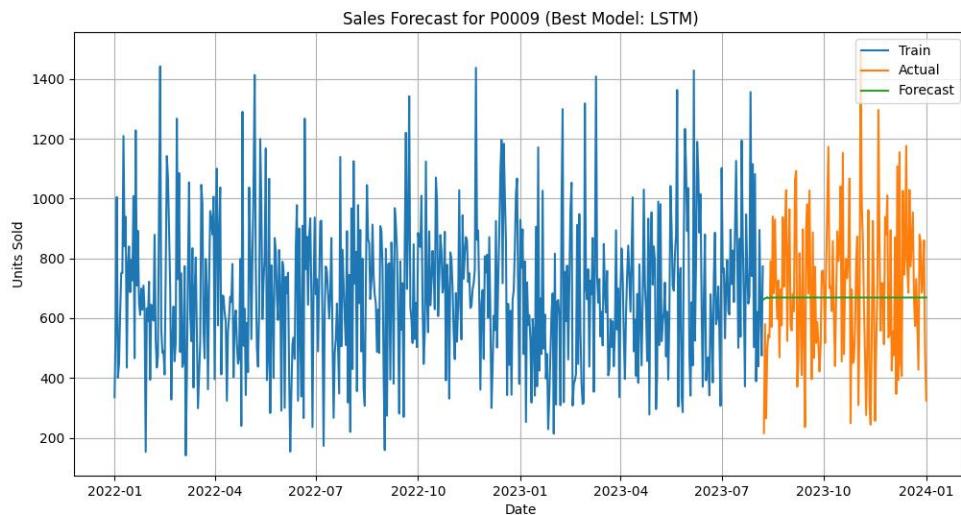
59 P0020 LSTM 176.975758 216.081290

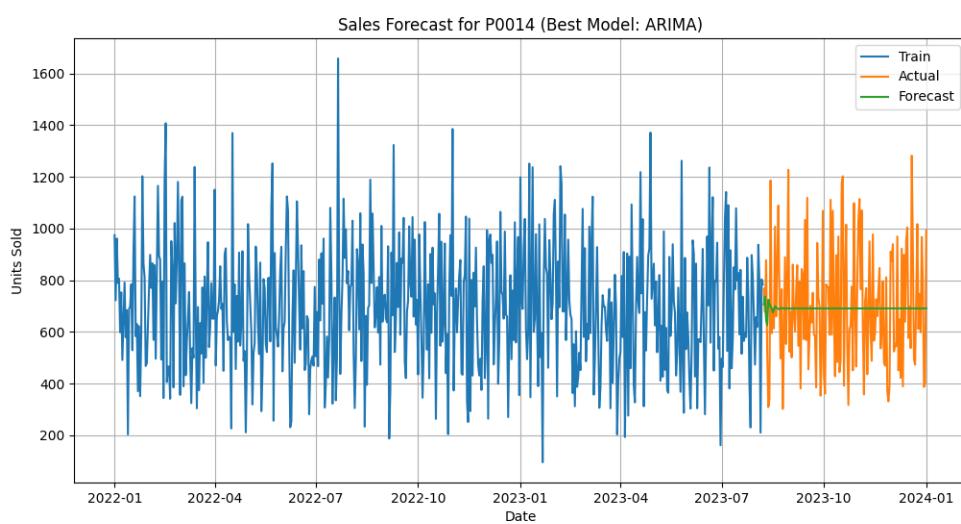
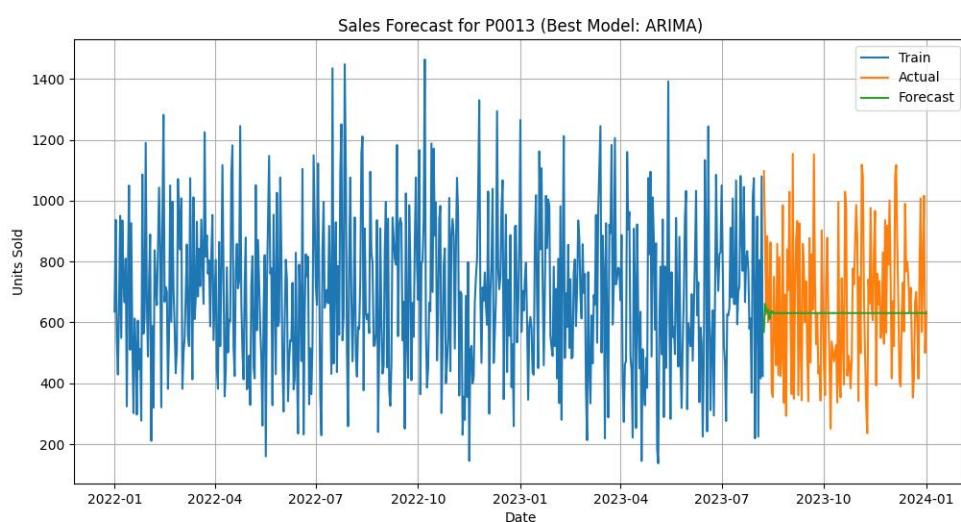
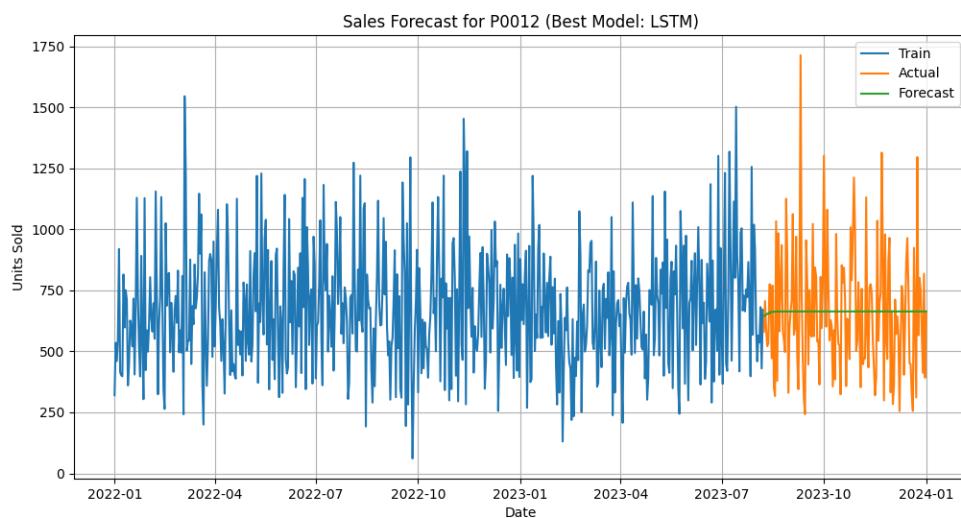
Forecast vs Actual Graphs for Each Product:

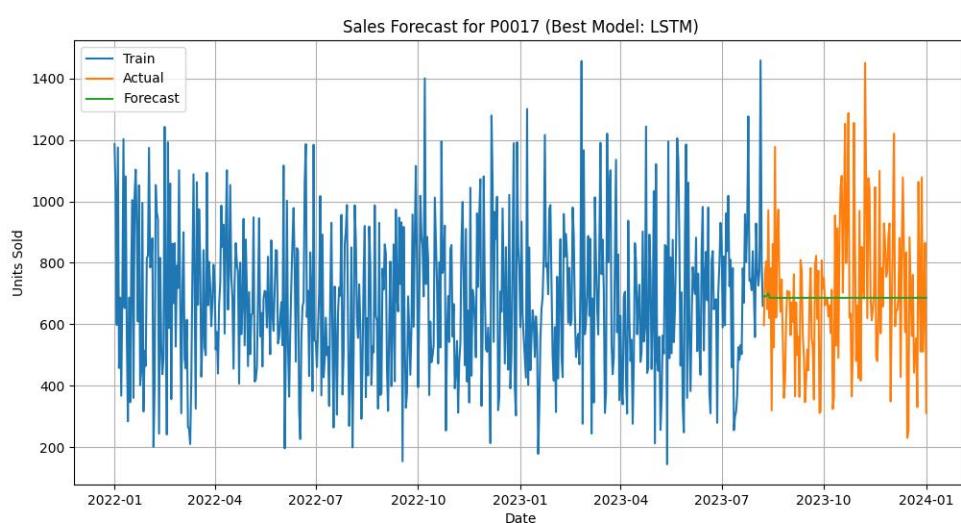
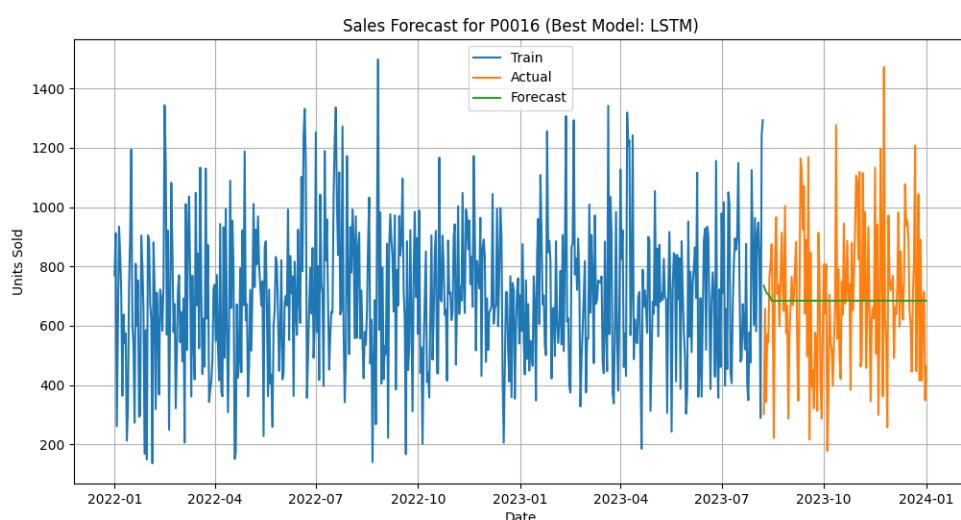
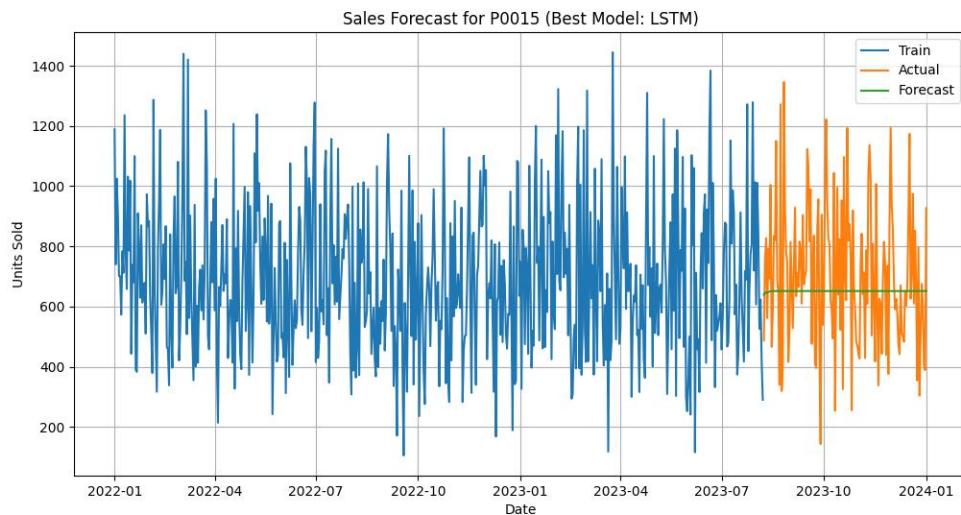


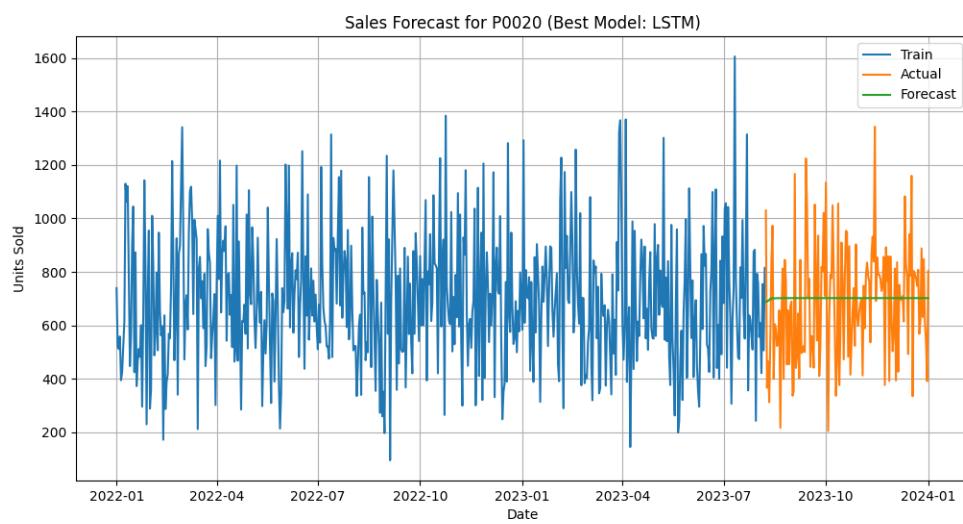
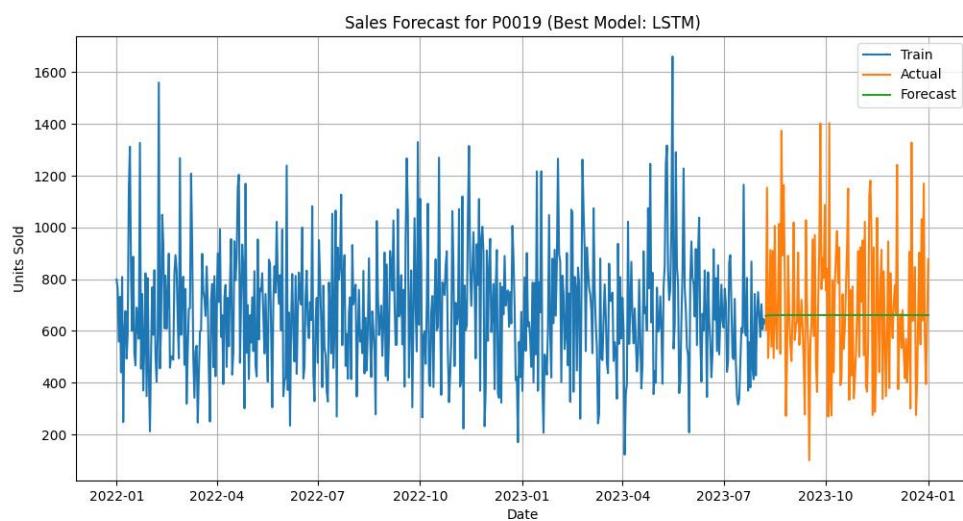
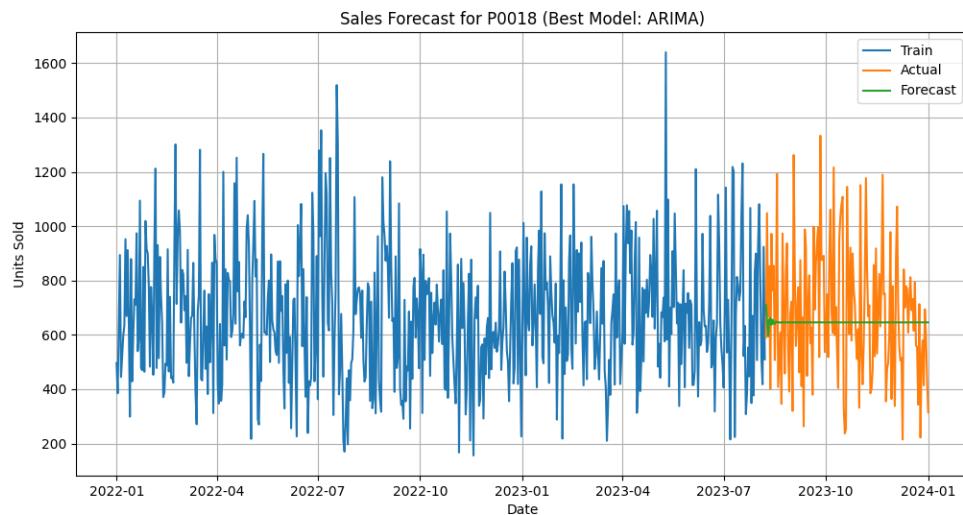


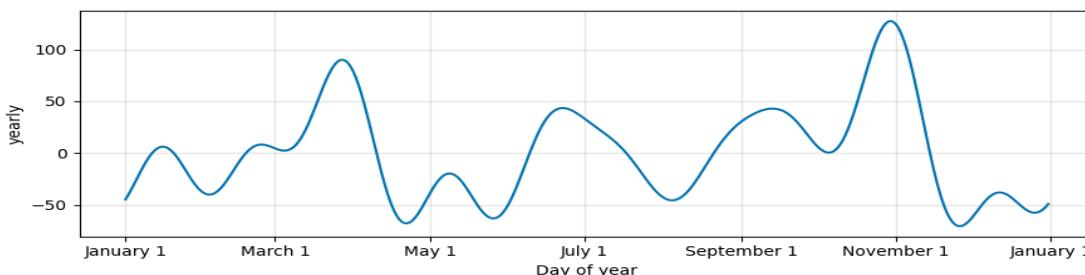
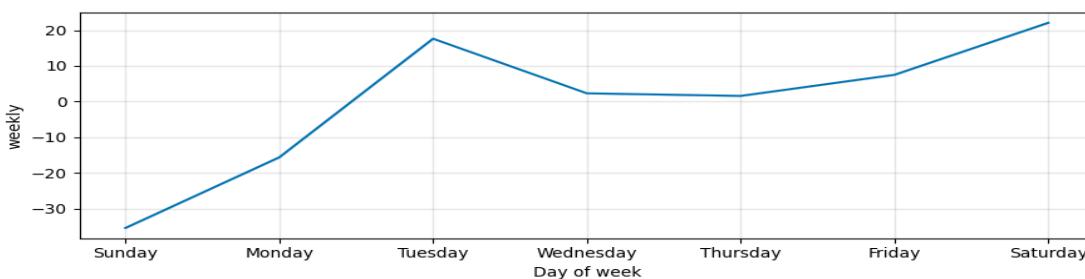
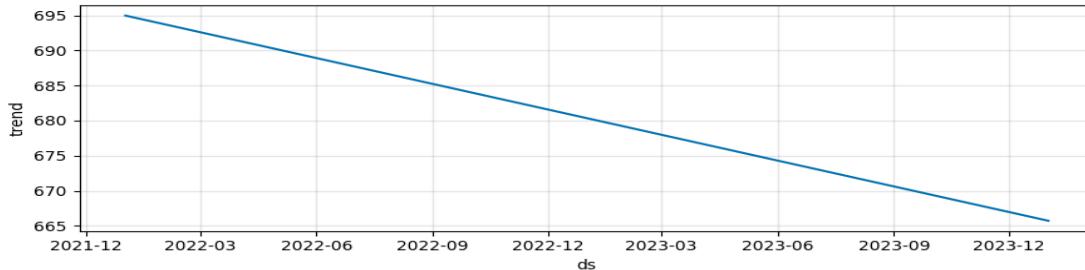
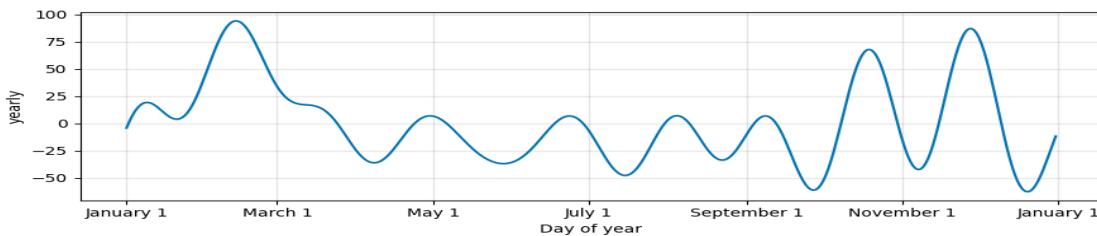
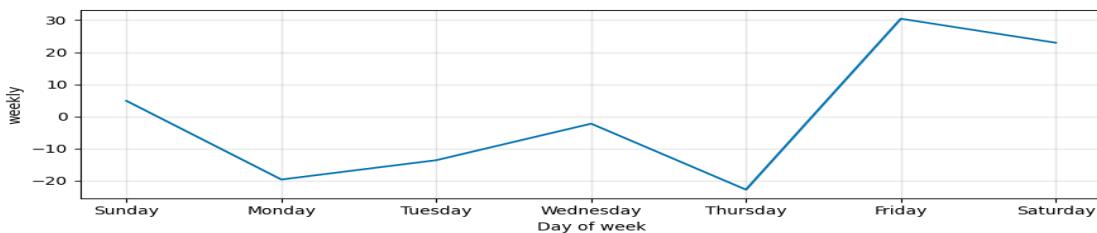
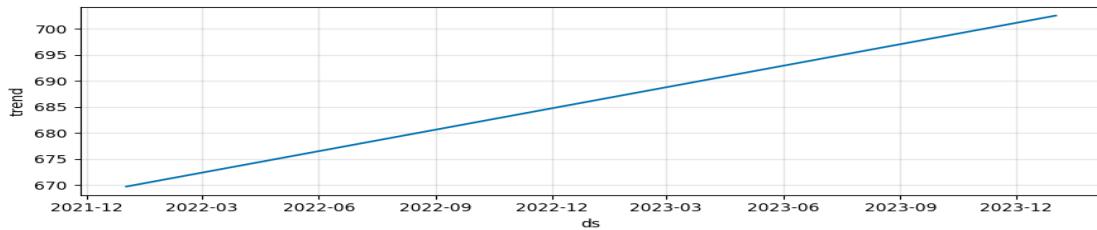


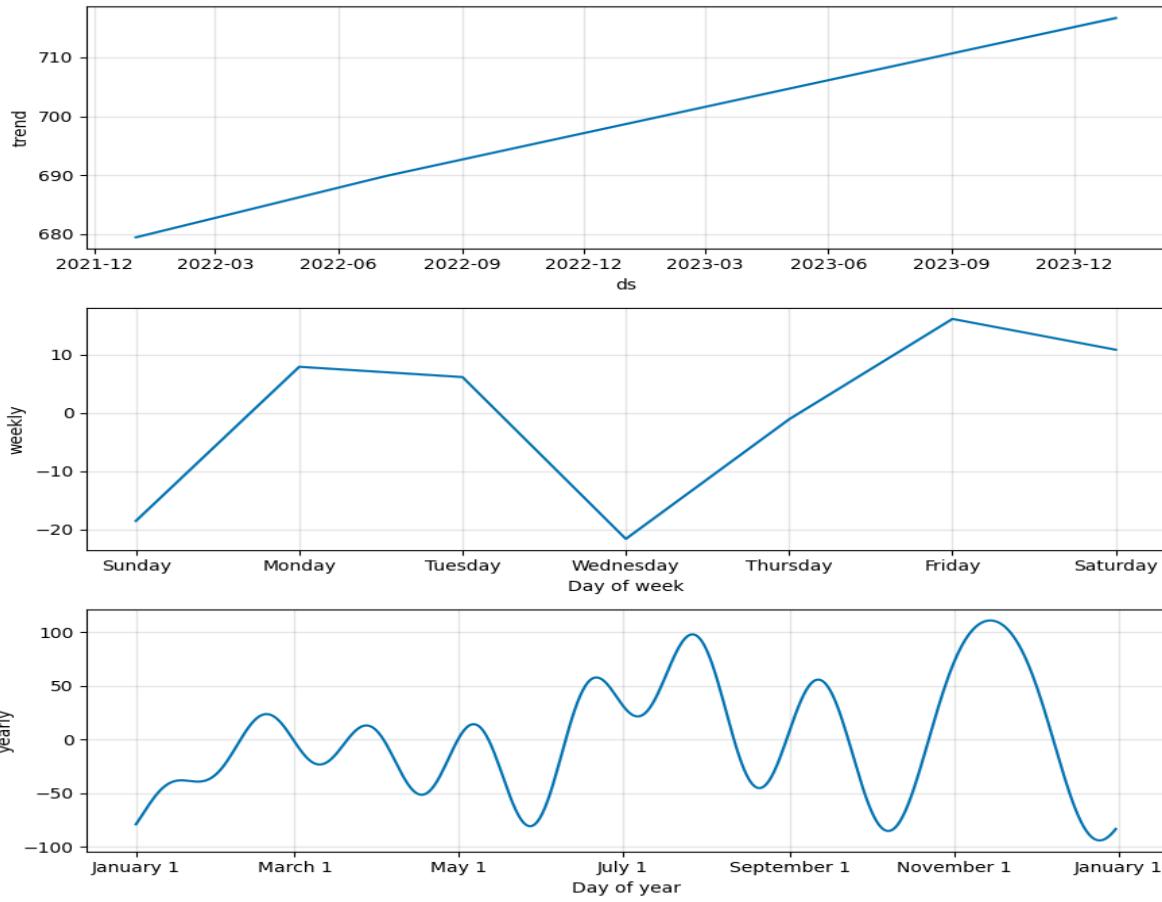












5. Observations

- Prophet captured seasonality well for products with strong periodic trends.
- ARIMA performed better for products with stable, linear demand patterns.
- LSTM models showed superior performance in handling nonlinear and complex temporal dependencies.
- Across products, LSTM and ARIMA were often selected as best models based on RMSE scores.

6. Outputs

The following key outputs were generated by the script:

- Best model for each product saved in the 'models/' directory.
- Forecast vs Actual plots saved in the 'plots/' directory.
- Evaluation metrics stored in 'data/evaluation_metrics.csv'.
- Forecasted results saved in 'data/forecast_results.csv'.

7. Conclusion & Next Steps

Milestone 2 successfully implemented forecasting models to predict future product sales trends. By evaluating Prophet, ARIMA, and LSTM models, the best model per product was identified and saved. These predictions will be used in the next milestone to optimize stock management and decision-making strategies based on future demand forecasts.